

# Scikit-learn (sklearn) Cheat Sheet

## 1. Library Structure Overview

**sklearn is modular and organized as follows:**

```
sklearn/
├── datasets          # Sample data, loading utilities
├── model_selection   # Splitting, cross-validation, hyperparam tuning
├── preprocessing     # Scaling, encoding, imputation, feature engineering
├── feature_selection # Selecting informative features
├── pipeline          # Chaining steps into workflows
├── metrics           # Scoring and evaluation functions
├── linear_model       # Linear/logistic regressions, elastic net, etc.
├── ensemble          # RandomForest, GradientBoosting, etc.
├── tree              # DecisionTreeClassifier/Regressor
├── neighbors         # k-NN algorithms
├── svm               # Support Vector Machines
├── naive_bayes       # Naive Bayes variants
├── cluster           # Clustering (KMeans, DBSCAN, etc.)
├── decomposition     # PCA, TruncatedSVD, etc.
├── manifold          # Advanced dimensionality reduction
├── compose           # ColumnTransformer, FeatureUnion
└── inspection        # Model explanation (permutation_importance, etc.)
```

## 2. Typical Machine Learning Workflow

### Step 1: Import Libraries

```
import numpy as np
import pandas as pd
from sklearn.model_selection import train_test_split
```

### Step 2: Load Data

```
from sklearn import datasets
X, y = datasets.load_iris(return_X_y=True)
```

Or from DataFrame:

```
df = pd.read_csv('file.csv')
X = df.drop('target', axis=1).values
y = df['target'].values
```

### Step 3: Split Data

```
X_train, X_test, y_train, y_test = train_test_split(
    X, y, test_size=0.2, random_state=42)
```

### Step 4: Preprocessing

```
from sklearn.preprocessing import StandardScaler, OneHotEncoder
scaler = StandardScaler().fit(X_train)
X_train_scaled = scaler.transform(X_train)
X_test_scaled = scaler.transform(X_test)
```

### Step 5: Build & Train Model

```
from sklearn.ensemble import RandomForestClassifier
clf = RandomForestClassifier(n_estimators=100, random_state=42)
clf.fit(X_train_scaled, y_train)
```

### Step 6: Predict & Evaluate

```
y_pred = clf.predict(X_test_scaled)
from sklearn.metrics import accuracy_score, classification_report, confusion_matrix
print("Accuracy:", accuracy_score(y_test, y_pred))
print(classification_report(y_test, y_pred))
```

## 3. Core Modules and Tasks

### A. Datasets

```
from sklearn.datasets import load_digits, fetch_openml
digits = load_digits()
```

```
# Or fetch larger datasets
X, y = fetch_openml('adult', as_frame=True, return_X_y=True)
```

## B. Preprocessing/Feature Engineering

```
from sklearn.preprocessing import MinMaxScaler, LabelEncoder

# Scaling
scaler = MinMaxScaler() # 0-1 scaling
X_scaled = scaler.fit_transform(X)

# Encoding
encoder = LabelEncoder()
y_encoded = encoder.fit_transform(y)
```

- **Pipelines** for sequential transformations:

```
from sklearn.pipeline import Pipeline
pipe = Pipeline([
    ('scaler', StandardScaler()),
    ('clf', RandomForestClassifier())
])
pipe.fit(X_train, y_train)
```

## C. Model Selection

Task	Function	Example
Train/test split	<code>train_test_split</code>	As above
Cross-validation	<code>cross_val_score</code>	<code>cross_val_score(clf, X, y, cv=5)</code>
Grid search	<code>GridSearchCV</code>	
Random search	<code>RandomizedSearchCV</code>	

### Example: Hyperparameter Search

```
from sklearn.model_selection import GridSearchCV
```

```
params = {'max_depth': [3, None], 'n_estimators': [10, 100, 200]}
gs = GridSearchCV(RandomForestClassifier(), params, cv=5)
gs.fit(X_train, y_train)
print("Best params:", gs.best_params_)
```

## D. Main Model Families

### Supervised Learning:

- **Classification:** LogisticRegression, RandomForestClassifier, SVC, GradientBoostingClassifier, KNeighborsClassifier, GaussianNB, MLPClassifier
- **Regression:** LinearRegression, Ridge, Lasso, SVR, RandomForestRegressor, GradientBoostingRegressor

### Unsupervised Learning:

- **Clustering:** KMeans, DBSCAN, AgglomerativeClustering, GaussianMixture
- **Dimensionality Reduction:** PCA, TruncatedSVD, NMF, TSNE, Isomap
- **Feature Selection:** SelectKBest, SelectFromModel, RFECV

## 4. Evaluation Metrics

### Classification:

```
from sklearn.metrics import accuracy_score, precision_score, recall_score, f1_score,
roc_auc_score
```

### Regression:

```
from sklearn.metrics import mean_absolute_error, mean_squared_error, r2_score
```

### Confusion Matrix and Reports:

```
from sklearn.metrics import confusion_matrix, classification_report
print(confusion_matrix(y_true, y_pred))
print(classification_report(y_true, y_pred))
```

## 5. Advanced Tools

## Pipelines & ColumnTransformer

- **Pipeline** for automated sequences:

```
from sklearn.pipeline import make_pipeline
pipe = make_pipeline(StandardScaler(), PCA(n_components=2), LogisticRegression())
```

- **ColumnTransformer** for feature-wise preprocessing:

```
from sklearn.compose import ColumnTransformer
from sklearn.impute import SimpleImputer
numeric_features = ['age', 'income']
categorical_features = ['gender', 'city']
ct = ColumnTransformer([
    ('num', StandardScaler(), numeric_features),
    ('cat', OneHotEncoder(), categorical_features)
])
X_trans = ct.fit_transform(df)
```

## Model Export & Persistence

```
from joblib import dump, load
dump(clf, 'rf_model.joblib')
clf_loaded = load('rf_model.joblib')
```

## Custom Transformers & Advanced Pipeline

```
from sklearn.base import BaseEstimator, TransformerMixin

class MyTransformer(BaseEstimator, TransformerMixin):
    def fit(self, X, y=None):
        # learn something from X
        return self
    def transform(self, X):
        # modify X
        return X
```

## Model Interpretation

```
from sklearn.inspection import permutation_importance
result = permutation_importance(clf, X_test, y_test)
importances = result.importances_mean
```

- Plot `feature_importances_` for tree-based models.

## 6. Special Topics and Pro Tips

- **Imbalanced Data:** `from imblearn.over_sampling import SMOTE`
- **Probabilities:** Use `predict_proba` for estimated class probabilities.
- **Chaining:** You can chain transformers and classifiers for end-to-end automation.
- **Parallelism:** Many models accept `n_jobs=-1` for multicore training/testing.

## 7. Example: End-to-End Workflow

```
import pandas as pd
from sklearn.model_selection import train_test_split, GridSearchCV
from sklearn.preprocessing import StandardScaler, OneHotEncoder
from sklearn.compose import ColumnTransformer
from sklearn.ensemble import GradientBoostingClassifier
from sklearn.pipeline import Pipeline

# Load data
df = pd.read_csv('titanic.csv')
y = df['Survived']
X = df.drop(['Survived', 'Name', 'Ticket'], axis=1)

# Preprocessing
num_features = ['Age', 'Fare']
cat_features = ['Sex', 'Embarked']

preprocessor = ColumnTransformer([
    ('num', StandardScaler(), num_features),
    ('cat', OneHotEncoder(), cat_features)
])

# Build pipeline
pipe = Pipeline([
```

```

    ('prep', preprocessor),
    ('clf', GradientBoostingClassifier())
])

# Train/test split
X_train, X_test, y_train, y_test = train_test_split(
    X, y, test_size=0.3, random_state=42)

# Hyperparameter tuning
params = {'clf__n_estimators': [100,200], 'clf__learning_rate': [0.01,0.1]}
gs = GridSearchCV(pipe, params, cv=5)
gs.fit(X_train, y_train)

print("Best parameters:", gs.best_params_)
print("Test accuracy:", gs.score(X_test, y_test))

```

## 8. Useful Links and References

- Official documentation/tutorials: [scikit-learn.org](https://scikit-learn.org)
- User Guide: Explanations and usage of each estimator and transformer
- Example Gallery: Dozens of annotated code examples
- Community cookbooks and notebooks

## 9. Summary Table: Common Scikit-learn Objects

Module	Key Classes/Functions	Usage
model_selection	<code>train_test_split</code> , <code>GridSearchCV</code> , <code>cross_val_score</code>	Splits, hyperparam search, CV
preprocessing	<code>StandardScaler</code> , <code>OneHotEncoder</code> , <code>LabelEncoder</code> , <code>SimpleImputer</code>	Prep/input wrangling
metrics	<code>accuracy_score</code> , <code>mean_squared_error</code> , <code>roc_auc_score</code>	Model evaluation
ensemble	<code>RandomForestClassifier</code> , <code>GradientBoostingClassifier</code>	Tree-based ensembles

pipeline	Pipeline, make_pipeline	Combine models & preprocessors
compose	ColumnTransformer	Parallel preprocessing per feature type
inspection	permutation_importance	Feature importances, explanation
decomposition	PCA, TruncatedSVD	Dimensionality reduction

## 10. Extra Advanced: Custom Scoring, Cross-validator, and Model Stacking

- **Custom scorer:**

```
from sklearn.metrics import make_scorer
scorer = make_scorer(f1_score, average='macro')
```

- **Nested cross-validation:**

```
from sklearn.model_selection import cross_val_score
cross_val_score(gs, X, y, cv=5)
```

- **Stack multiple models:**

```
from sklearn.ensemble import StackingClassifier
stack = StackingClassifier(
    estimators=[('lr', LogisticRegression()), ('rf', RandomForestClassifier())],
    final_estimator=GradientBoostingClassifier()
)
stack.fit(X_train, y_train)
```